

ECE586 MP1 Markov Chains

Exercise 1

What is the distribution of the number of fair coin tosses before one observes 3 heads in a row? To solve this, consider a 4-state Markov chain with transition probability matrix

$$P = \begin{bmatrix} 0.5 & 0.5 & 0 & 0 \\ 0.5 & 0 & 0.5 & 0 \\ 0.5 & 0 & 0 & 0.5 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

where $X_t = 1$ if the previous toss was tails, $X_t = 2$ if the last two tosses were tails then heads, $X_t = 3$ if the last three tosses were tails then heads twice, and $X_t = 4$ is an absorbing state that is reached when the last three tosses are heads.

- **(15 pts)** Write a computer program (e.g., in Python, Matlab, ...) to compute $\Pr(T_{1,4} = m)$ for $m = 1, 2, \dots, 100$ and **(10 pts)** use this to compute and print an estimate of the expected number of tosses $E[T_{1,4}]$


```
In [4]: # See compute_Phi_ET in fsmc_code.py

P = np.array([[0.5, 0.5, 0, 0], [0.5, 0, 0.5, 0], [0.5, 0, 0, 0.5], [0, 0, 0, 1]])
Phi_list, ET = compute_Phi_ET(P, 100)

m = np.arange(1, 101)                                     ### steps to be plotted
Pr = Phi_list[m, 0, 3] - Phi_list[(m - 1), 0, 3]          ### \Pr(T_{1,4} = m) for all m
E = ET[0, 3]                                                ### \mathbb{E}[T_{1,4}]

plt.figure(figsize=(10, 6))
plt.stem(m, Pr, use_line_collection=True)
plt.xlabel(r'm')
plt.ylabel(r'\Pr(T_{1,4}=m)')
plt.title(r'E[T_{1,4}] = ' + str(E))
```

Out[4]: Text(0.5, 1.0, 'E[T_{1,4}] = 13.970568255020394')




- **(15 pts)** Write a computer program that generates 500 realizations from this Markov chain and **(10 pts)** uses them to plot a histogram of $T_{1,4}$ and compute/print an estimate of the expected number of tosses $\mathbb{E}[T_{1,4}]$.

In [5]: `# implement simulate_hitting_time(P, states, nr) in fsmc_code.py`

```
T = simulate_hitting_time(P, [0, 3], 500)
plt.figure(figsize=(10, 6))
plt.hist(T, bins=np.arange(max(T))+0.5)
plt.title(r'mean of  $T_{\{1,4\}}$  = ' + str(np.mean(T)))
```

Out[5]: Text(0.5, 1.0, 'mean of $T_{\{1,4\}}$ = 14.276')



Exercise 2

Consider the miniature chutes and ladders game shown in Figure 1. Assume a player starts on the space labeled 1 and plays by rolling a fair four-sided die and then moves that number of spaces. If a player lands on the bottom of a ladder, then they automatically climb to the top. If a player lands at the top of a slide, then they automatically slide to the bottom. This process can be modeled by a Markov chain with $n = 16$ states where each state is associated with a square where players can start their turn (e.g., players never start at the bottom of a ladder or the top of a slide). To finish the game, players must land exactly on space 20 (moves beyond this are not taken).

- **(10 pts)** Compute the transition probability matrix P of the implied Markov chain.

```
In [6]: # You can either do this by hand (e.g., look at picture and write down matrix) or by automating the process.
```

```
# By hand
# P = np.asarray([[...], [...], [...], ...])

###

# Or automated general function for Chutes and Ladders games
def construct_P_matrix(n, dice, chutes, ladders):
    """
    Arguments:
        n (int) -- size of the state space
        dice (numpy.array) -- probability distribution of the dice outcome
        chutes (list[(int, int)]) -- the list of chutes, in pairs of (start, end)
        ladders (list[(int, int)]) -- the list of ladders, in pairs of (start, end)

    Returns:
        P (numpy.array) -- n x n, transition matrix of the Markov chain
    """
    # Add code here to build matrix

    #####
    P = np.zeros((n, n))
    num_ladders = len(ladders)
    num_chutes = len(chutes)
    dice_length = len(dice)
    n_expand = n + num_ladders + num_chutes
    P_expand = np.zeros((n_expand, n_expand))
    remove = [0] * (num_ladders + num_chutes)

    for i in range(0, n_expand):
        if i + dice_length < n_expand:
            P_expand[i, (i+1):(i+1+dice_length)] = dice
        else:
            P_expand[i, i] = 1 - sum(dice[0:(n_expand-(i+1))])
            P_expand[i, (i+1):n_expand] = dice[0:(n_expand-(i+1))]

    for j in range(0, num_ladders):
        P_expand[:, (ladders[j][1]-1)] = P_expand[:, (ladders[j][1]-1)] + P_expand[:, (ladders[j][0]-1)]
        remove[j] = ladders[j][0] - 1

    for k in range(0, num_chutes):
        P_expand[:, (chutes[k][1]-1)] = P_expand[:, (chutes[k][1]-1)] + P_expand[:, (chutes[k][0]-1)]
        remove[num_ladders + k] = chutes[k][0] - 1


    P_remove_row = np.delete(P_expand, remove, 0)
    P_remove_row_column = np.delete(P_remove_row, remove, 1)
    P = P_remove_row_column
    P[n-1, n-1] = 1
    #####

    return P

n = 16 ### number of states
dice = np.array([1/4, 1/4, 1/4, 1/4]) ### probability distribution of dice
chutes = [(13, 2), (17, 6)] ### (source, destination) pairs of chutes
ladders = [(4, 8), (14, 19)] ### (source, destination) pairs of ladders
P = construct_P_matrix(n, dice, chutes, ladders)
###

# Plot transition matrix
plt.figure(figsize=(8, 8))
plt.imshow(P)
```

```
Out[6]: <matplotlib.image.AxesImage at 0x28764411100>
```




- **(10 pts)** For this Markov chain, use function from Exercise 1 to compute and plot the cumulative distribution of the number turns a player takes to finish (i.e., the probability $\Pr(T_{1,20} \leq m)$ where $T_{1,20}$ is the hitting time from state 1 to state 20). Compute and print the mean $E[T_{1,20}]$.

```
In [7]: # Use previous functions to complete this exercise
Phi_list, ET = compute_Phi_ET(P, ns=100)

m = np.arange(1, 101)  ### steps to be plotted
Pr = Phi_list[m, 0, 15]  ### \Pr(T_{1,20} <= m) for all m
E = ET[0, 15]  ### \mathbb{E}[T_{1,20}]

plt.figure(figsize=(10, 6))
plt.plot(m, Pr)
plt.xlabel(r'm')
plt.ylabel(r'\Pr(T_{1,20} <= m)')
plt.title(r'E[T_{1,20}] = ' + str(E))
```


Out[7]: Text(0.5, 1.0, 'E[T_{1,20}] = 12.792331782072552')



- **(10 pts)** Use function from Exercise 1 to generate 500 realizations from this Markov chain. Then, use them to plot a histogram of $T_{1,20}$ and compute/print an estimate of the expected number of tosses $\mathbb{E}[T_{1,20}]$.

```
In [8]: # Use previous functions to complete this exercise
T = simulate_hitting_time(P, [0, n-1], 500)
plt.figure(figsize=(10, 6))
plt.hist(T, bins=np.arange(max(T))-0.5)
plt.title(r'mean of  $T_{1,20}$  =' + str(np.mean(T)))
```

```
Out[8]: Text(0.5, 1.0, 'mean of  $T_{1,20}$  =12.598')
```



Exercise 3

- (10 pts)** Write a program to compute the stationary distribution of a Markov chain when it is unique. Consider a game where the gameboard has 8 different spaces arranged in a circle. During each turn, a player rolls two 4-sided dice and moves clockwise by a number of spaces equal to their sum. **(5 pts)** Define the transition matrix for this 8-state Markov chain and compute its stationary distribution.

```
In [11]: # Use previous functions to complete this exercise
### construct the transition matrix
#####
P = np.zeros((8, 8))
### [1/16, 0, 1/16, 2/16, 3/16, 4/16, 3/16, 2/16]
P[0, :] = [1/16, 0, 1/16, 1/8, 3/16, 1/4, 3/16, 1/8]
for row_state in range(1, 8) :
    P[row_state, :] = np.roll(P[row_state-1, :], 1)
#####
print(P)
stationary_distribution(P)
```

```
[[0.0625 0.         0.0625 0.125  0.1875 0.25   0.1875 0.125 ]
 [0.125  0.0625 0.         0.0625 0.125  0.1875 0.25   0.1875]
 [0.1875 0.125  0.0625 0.         0.0625 0.125  0.1875 0.25  ]
 [0.25   0.1875 0.125  0.0625 0.         0.0625 0.125  0.1875]
 [0.1875 0.25   0.1875 0.125  0.0625 0.         0.0625 0.125 ]
 [0.125  0.1875 0.25   0.1875 0.125  0.0625 0.         0.0625]
 [0.0625 0.125  0.1875 0.25   0.1875 0.125  0.0625 0.        ]
 [0.         0.0625 0.125  0.1875 0.25   0.1875 0.125  0.0625]]
```

```
Out[11]: array([[0.125],
 [0.125],
 [0.125],
 [0.125],
 [0.125],
 [0.125],
 [0.125],
 [0.125]])
```

Next, suppose that one space is special (e.g., state-1 of the Markov chain) and a player can only leave this space by rolling doubles (i.e., when both dice show the same value). Again, the player moves clockwise by a number of spaces equal to their sum. **** (5 pts) **** Define the transition matrix for this 8-state Markov chain and compute its stationary probability distribution.

```
In [12]: # Use previous functions to complete this exercise
        ### construct the transition matrix
        #####
        P = np.zeros((8, 8))
        P[0, :] = [13/16, 0, 1/16, 0, 1/16, 0, 1/16, 0]
        ###          [1/16, 0, 1/16, 2/16, 3/16, 4/16, 3/16, 2/16]
        P[1, :] = np.roll([1/16, 0, 1/16, 1/8, 3/16, 1/4, 3/16, 1/8], 1)
        for row_state in range(2, 8) :
            P[row_state, :] = np.roll(P[row_state-1, :], 1)
        #####
        print(P)
        stationary_distribution(P)

[[0.8125 0.      0.0625 0.      0.0625 0.      0.0625 0.      ]
 [0.125  0.0625 0.      0.0625 0.125  0.1875 0.25   0.1875]
 [0.1875 0.125  0.0625 0.      0.0625 0.125  0.1875 0.25   ]
 [0.25   0.1875 0.125  0.0625 0.      0.0625 0.125  0.1875]
 [0.1875 0.25   0.1875 0.125  0.0625 0.      0.0625 0.125  ]
 [0.125  0.1875 0.25   0.1875 0.125  0.0625 0.      0.0625]
 [0.0625 0.125  0.1875 0.25   0.1875 0.125  0.0625 0.      ]
 [0.      0.0625 0.125  0.1875 0.25   0.1875 0.125  0.0625]]
```

```
Out[12]: array([[0.41836864],
                [0.08285234],
                [0.10176963],
                [0.07092795],
                [0.09311176],
                [0.0625555 ],
                [0.09593429],
                [0.07447989]])
```

In []: